

Project 2: Human Detection

Abhedya N Khatiwala

University ID: N19478833

Net ID: ank434

Email ID: ank434@nyu.edu

Course: CS6643 Computer Vision and Scene Analysis

a) File Name of the source code:

The name of the python executable file is “HumanDetection.py”

The files that contain the HOG Descriptors for “crop001045b.bmp” and “crop001278a.bmp” are labelled “crop001045b.txt” and “crop001278a.txt” respectively.

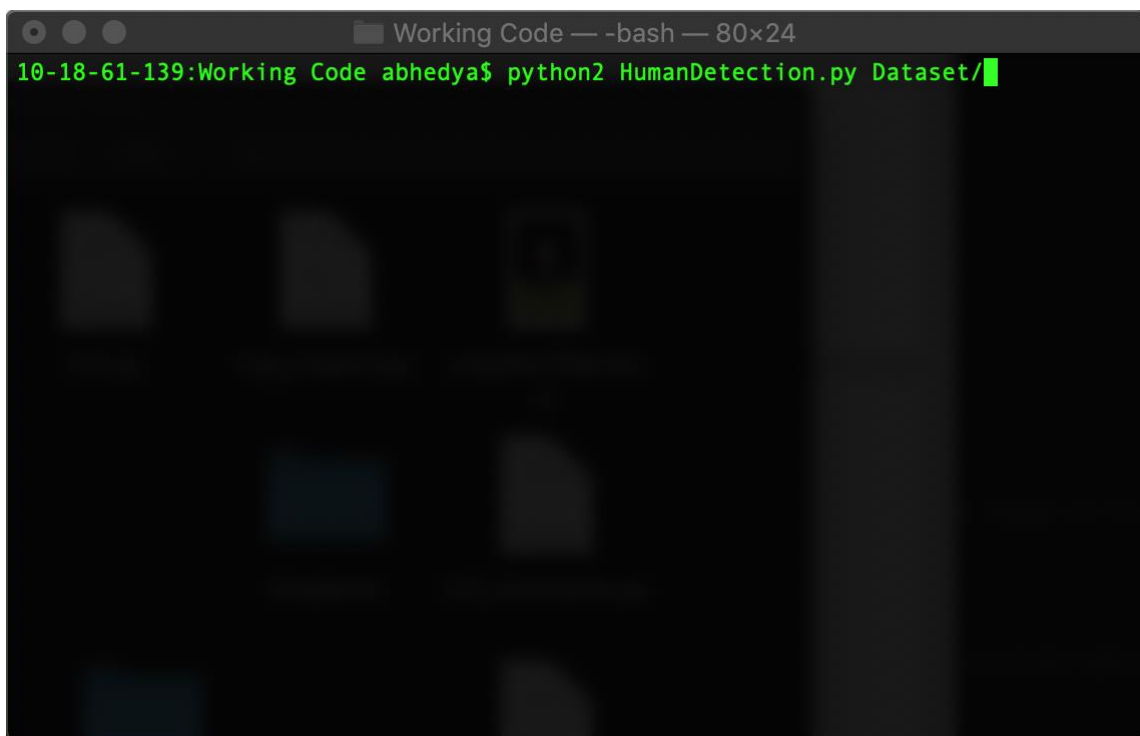
b) Instructions for running the program:

Initially, begin by pasting the folder containing the Dataset in the same folder as the python code. This is just for ease of executing the program.

Please note that, the Dataset folder must contain the Folders “Test_Neg”, “Test_Positive”, “Train_Positive”, “Train_Negative” with their correct respective images inside them.

For executing the code, type the following command on the terminal:

python2 HumanDetection.py Dataset/

A screenshot of a terminal window titled "Working Code — -bash — 80x24". The terminal shows the command "python2 HumanDetection.py Dataset/" being executed. The prompt is "10-18-61-139:Working Code abhedya\$". The background of the terminal is dark with some faint, blurry icons visible.

```
10-18-61-139:Working Code abhedya$ python2 HumanDetection.py Dataset/
```

Required Libraries: math, numpy, matplotlib, cv2, sys, os, glob

c)

Q1) How did you initialize the weight values of the network?

Ans1) I randomly initialized the weights using the `np.random.rand()` function, which generated random numbers.

Q2) How many iterations (or epochs) through the training data did you perform?

Ans2) I performed 100 epochs.

Q3) How did you decide when to stop training?

Ans3) When the weights started to converge and the error between the actual value and predicted value became very small.

Q4) Based on the output value of the output neuron, how did you decide on how to classify the input image into human or not-human?

Ans4) If the predicted probability for a certain image was above 50%, the neural network classified the image as containing a Human. If the probability was below 50% the neural network classified the image as not containing a Human.

d and e) Following are the results achieved. Note the parameters mentioned below every image table. The values in the Column “Output Value” are the probabilities predicted by the neural network.

Test Image	Output value	Classification
crop_000010b	0.92160128	Human
crop001008b	0.60681299	Human
crop001028a	0.58773214	Human
crop001045b	0.87280838	Human
crop001047b	0.77229701	Human
00000053a_cut	0.09512104	Not Human
00000062a_cut	0.32254887	Not Human
00000093a_cut	0.32127101	Not Human
no_person__no_bike_213_cut	0.2817215	Not Human
no_person__no_bike_247_cut	0.2429066	Not Human

■ Epochs: 100, Hidden Neurons: 1000

Test Image	Output value	Classification
crop_000010b	0.90775864	Human
crop001008b	0.62555973	Human
crop001028a	0.57816784	Human
crop001045b	0.90627158	Human
crop001047b	0.79141005	Human
00000053a_cut	0.11248245	Not Human
00000062a_cut	0.30224689	Not Human
00000093a_cut	0.33281027	Not Human
no_person__no_bike_213_cut	0.24704088	Not Human
no_person__no_bike_247_cut	0.22784613	Not Human



■ Epochs: 100, Hidden Neurons: 500

Test Image	Output value	Classification
crop_000010b	0.92358462	Human
crop001008b	0.64321735	Human
crop001028a	0.57564167	Human
crop001045b	0.90231501	Human
crop001047b	0.78677533	Human
00000053a_cut	0.08350137	Not Human
00000062a_cut	0.27950736	Not Human
00000093a_cut	0.36318688	Not Human
no_person__no_bike_213_cut	0.24066211	Not Human
no_person__no_bike_247_cut	0.23919398	Not Human


■ Epochs: 100, Hidden Neurons: 250

f) Normalized Gradient Magnitude Images for all the test images

Positive Test Images



Negative Test Images



g) Source Code

#The author of this code is "Abhedya N Khatiwala"

#CS Grad at NYU School of Engineering

```
import math
import numpy as np
from PIL import Image
from sys import argv
import cv2
import matplotlib.pyplot as plt
from matplotlib.pyplot import imshow, show, subplot, figure, gray, title, axis, savefig, imread
import sys
import os
import glob
```

```
class HistogramOfGradients:
```

```
    def findImages(self, dataFile, delimiter): #This function takes data folders and returns a
list of images and another list which contains actual output for that image
```

```
        PathList = []
```

```
        dataOut = []
```

```
        for dataFolder in dataFile.keys():
```

```
            for directoryName, subDirectory, fileL in os.walk(dataFolder):
```

```
                for imageFile in fileL:
```

```
                    imageP = dataFolder + delimiter + imageFile
```

```
                    PathList.append(imageP)
```

```
                    dataOut.append([dataFile[dataFolder]])
```

```
        return PathList, dataOut
```

```
    def RGBtoGray(self, image):    #this function converts the given RGB image in to a
Grayscale Image
```

```
        return np.round(np.dot(image[...,:3], [0.299, 0.587, 0.114]))
```

```

def convolveP(self, image, kernel):    #this function performs the convolution
operation.

    iN = len(image)
    jN = len(image[0])
    image_new = np.zeros((iN, jN), dtype = np.float) #initializing a matrix with
float(0)
    for a in range(0, iN - 2):
        for b in range(0, jN - 2):
            image_new[a+1][b+1] = (np.sum(image[a: a + 3,
convolution
                                   b: b + 3] * kernel))/3 #multiplication of two matrices for

    return image_new


def prewittGradient(self, image):

    i = len(image)
    j = len(image[0])

    op1 = np.array([-1, 0, 1],
                    [-1, 0, 1],
                    [-1, 0, 1])) #Prewitt's Operator

    op2 = np.array([1, 1, 1],
                    [0, 0, 0],
                    [-1, -1, -1])) #Prewitt's Operator

    Gx = self.convolveP(image, op1)
    #Gx = Gx/3
    Gy = self.convolveP(image, op2)
    #Gy/3
    gradient = np.zeros((i, j), dtype = np.float)

    grad_angle = np.zeros((i, j), dtype = np.float)

    for m in range(i):
        for n in range(j):
            gradient[m][n] = (np.sqrt(np.square(Gx[m][n]) +
np.square(Gy[m][n])))/np.sqrt(2) # computing gradient for every pixel value and normalizing it

```

angles

```
#here we compute the grad angle by and check for negative

grad_angle[m][n] = 0
if(Gx[m][n] == 0):
    if(Gy[m][n] == 0):
        grad_angle[m][n] = 0
    # grad_angle[m][n] = 0
    else:
        if(Gy[m][n] > 0):
            grad_angle[m][n] = 90
        else:
            grad_angle[m][n] = -90

else:
    grad_angle[m][n] =
math.degrees(np.arctan(Gy[m][n]/Gx[m][n])) #computing gradient angle for each pixel

if(grad_angle[m][n] < 0):
    grad_angle[m][n] = grad_angle[m][n] + 180

return gradient, grad_angle
```

def histogram(self, grad_angle, gradient): #this function computes the histogram and returns the descriptor.

```
i = len(gradient)
j = len(gradient[0])

featureVec = []
cell_size = (8, 8)    #defining the cell size
block_size = (16, 16) #defining the block size
step_size = (8, 8)    #defining the bin size
grad_angle = ((grad_angle)/20.0)    #calculating the the

i_cell_size = cell_size[0]    #cell size assigned for row
j_cell_size = cell_size[1]    #cell size assigned for column

i_cell_count = i//cell_size[0]    #calculating the cell count
j_cell_count = j//cell_size[1]
```

```
i_cells_per_block, j_cells_per_block = np.array(block_size)/np.array(cell_size)
#calculating the number of cells per block
```

```
i_cells_per_step, j_cells_per_step = np.array(step_size)/np.array(step_size)
```

```
i_block_count = (i_cell_count - i_cells_per_block) / i_cells_per_step + 1
j_block_count = (j_cell_count - j_cells_per_block) / j_cells_per_step + 1
```

```
bin_histograms = np.zeros((i_cell_count, j_cell_count, 9)) #we create an array
to store the histogram of all cells.
```

```
#this fills the histogram array by going through the gradient angles.
for I in range(i):
    for J in range(j):
```

```
current pixel          c_angle = grad_angle[I][J]    #calculating the edge magnitude at
```

```
current pixel          c_mag = gradient[I][J]      #calculating the gradient angle at
```

```
l_bin = np.floor(c_angle)    #calculating the bins
h_bin = np.ceil(c_angle)
```

```
nearby bins           l_dist = c_angle - l_bin    #calculating the angle distance to it's
h_dist = c_angle - h_bin
```

```
from the center of two bins.
l_val = c_mag * h_dist      #calculating the distance based on
h_val = c_mag * l_dist
```

```
bin_histograms[int(I // i_cell_size)][int(J / j_cell_size)][int(l_bin)]
+= l_val
```

```
bin_histograms[int(I // i_cell_size)][int(J /
j_cell_size)][int(h_bin%9)] += h_val
```

```
bin_histograms = bin_histograms.astype(np.float)
```

```
#here we perform Block Level Normalization
```

```

    for iWc in range(0, i_block_count, i_cells_per_step):
        for jWc in range(0, j_block_count, j_cells_per_step):

            c_block = bin_histograms[iWc: iWc + i_cells_per_block, jWc: jWc +
j_cells_per_block].reshape(-1)
            c_block_l2 = sum(c_block**2)**0.5

            if c_block_l2 != 0:
                c_block /= c_block_l2
                featureVec.append(abs(c_block))

featureVec = np.array(featureVec).reshape(-1, 1)

return featureVec

```

#This function is like a main function in this Class. It takes as input, a list which contains images and gives back

#an array of feature vectors.

```
def hog(self, im_path):
```

```
    features = []
```

```
    for images in im_path:
```

```
        inputImage = self.GetImages(images)
```

```
        if('png' in sys.argv[1]):
```

```
            inputImage *= 255
```

Grayscale image

```
        grayscale = self.RGBtoGray(inputImage) #this function call returns a
```

```
        #gsResult = Image.fromarray(grayscale)
```

and the angle and using Prewitt's Operator

```
        gradient, angle = self.prewittGradient(grayscale) #computing the gradient
```

```
        #gradientResult = Image.fromarray(gradient)
```

```
        cellRes = self.histogram(angle, gradient)
```

```
        features.append(cellRes)
```

```
    return features
```

```
def GetImages(self, im_path):    #this function reads image from a file
    return np.array(plt.imread(im_path))
```

#This class consists of all the functions required to initialize, train and test the Neural Network.

```
class NeuralNetwork:
```

```
    #initializing the neural net parameters
```

```
    # arch = (7524, 250, 1) indicates the input neurons, hidden neurons, the epochs and the
learning_rate
```

```
    # for changing the number of hidden neurons, change the value 250/500/1000
```

```
    def __init__(self, arch = (7524, 250, 1), epochs = 100, learning_rate = 0.01):
```

```
        self.arch = arch
```

```
        #we initialize the weights with Random Values
```

```
        self.weights1 = np.random.randn(arch[1], arch[0]) * 0.01
```

```
        self.weights2 = np.random.randn(arch[2], arch[1]) * 0.01
```

```
        #we initialize the bias with zeroes
```

```
        self.bias1 = np.zeros((arch[1], 1))
```

```
        self.bias2 = np.zeros((arch[2], 1))
```

```
        self.layer1 = self.layer2 = None
```

```
        self.d_weights1 = self.d_weights2 = None
```

```
        self.epochs = epochs
```

```
        self.learning_rate = learning_rate
```

```
#this function performs the forward propogation.
```

```
def feedforward(self, training_data):
```

```
    a1 = self.weights1.dot(training_data) + self.bias1
```

```
    self.layer1 = self.ReLU(a1)
```

```
    self.layer2 = self.sigmoid(self.weights2.dot(self.layer1) + self.bias2)
```

```
#this function calculates the error on every epoch
```

```
def error(self, actual_output):
```

```
    return 0.5*np.square(self.layer2 - actual_output).sum()
```

#this function performs the back propogation after the error is calculated and rectifies the weights

```
def backprop(self, training_data, actual_output):
    diff = self.layer2 - actual_output
    z2 = diff * self.sigmoid_derivative(self.layer2)
    self.d_weights2 = np.dot(z2, self.layer1.T)

    z1 = np.dot(self.weights2.T, z2) * self.dReLU(self.layer1)
    self.d_weights1 = np.dot(z1, training_data.T)

    self.d_bias2 = np.sum(z2, axis = 1, keepdims = True)
    self.d_bias1 = np.sum(z1, axis = 1, keepdims = True)
```

#this function correctly updates the weights.

```
def update(self):
    self.weights1 = self.weights1 - self.learning_rate * self.d_weights1
    self.bias1 = self.bias1 - self.learning_rate * self.d_bias1

    self.weights2 = self.weights2 - self.learning_rate * self.d_weights2
    self.bias2 = self.bias2 - self.learning_rate * self.d_bias2
```

#this function trains the neural net and invokes the self.function defined in this class

```
def trainNN(self, trainImages, training_data_in, training_data_out):
    trainLen = len(training_data_in)
    for epoch in range(self.epochs):
        ep_error = 0.0 #we initialize the epoch_error initially as zero
        for data_count, train_data in enumerate(training_data_in):
            self.feedforward(train_data) #for ever image in the
dataset, we call the feedforward function
            error = self.error(training_data_out[data_count]) #we
compute the error by calling the error function defined in this class
            ep_error += error #we update the error
            self.backprop(train_data, training_data_out[data_count]) #then
we call the backprop function defined in this class to update the weights
            self.update() #this function correctly updates the weights that
we figured out using backpropogation
```

```

print("Epoch Count: " + str(epoch), "Average Error: ", ep_error/trainLen)

#this function tests the neural net and returns the misclassified count.
#this function basically predicts the probability of whether a human is present in an
image.
#if the probability is above 0.5, then it classifies the image as human.
def testNN(self, testImages, testing_data_in, testing_data_out):
    misclassify = 0

    positiveList = []
    negativeList = []

    for data_count, test_data in enumerate(testing_data_in):
        self.feedforward(test_data)
        print("Predicted Probability: " + str(self.layer2), "Actual Probability Value:
" + str(testing_data_out[data_count]))

        cPrediction = np.round(self.layer2.sum())

        if cPrediction:
            positiveList.append([testImages[data_count],
str(self.layer2.sum())])
        else:
            negativeList.append([testImages[data_count],
str(self.layer2.sum())])

        misclassify += (float(cPrediction - testing_data_out[data_count]) == 0)

    #Fig = plt.figure()

    print(str(float(misclassify) / float(len(testing_data_out)) * 100) + " % Prediction
Accuracy")
    #print(len(testing_data_out))

    #sigmoid function
    def sigmoid(self, t):
        return 1/(1+np.exp(-t))

    #this function calculates the sigmoid derivative required for backpropagation
    def sigmoid_derivative(self, p):

```



```
return p * (1 - p)
```

```
#this function is the RELU activation function
```

```
def ReLU(self, t):
```

```
    return np.maximum(t, 0)
```

```
#this function is the derivative of ReLU used for backpropogation
```

```
def dReLU(self, t):
```

```
    return 1*(t>0)
```

#The following code basically extracts the images present in the Dataset folder, and calls the respective functions.

```
if __name__ == "__main__":
```

```
    if len(sys.argv) >= 2:
```

```
        dataPath = sys.argv[1]
```

```
    else:
```

```
        dataPath = raw_input("No Path to Data Found. Enter the path to the data
```

```
directory: ")
```

```
positiveTrainDataFolder = 'Train_Positive'
```

```
negativeTrainDataFolder = 'Train_Negative'
```

```
positiveTestDataFolder = 'Test_Positive'
```

```
negativeTestDataFolder = 'Test_Neg'
```

```
delimiter = '/'
```

```
positiveTrainPath = dataPath + delimiter + positiveTrainDataFolder
```

```
negativeTrainPath = dataPath + delimiter + negativeTrainDataFolder
```

```
positiveTestPath = dataPath + delimiter + positiveTestDataFolder
```

```
negativeTestPath = dataPath + delimiter + negativeTestDataFolder
```

```
trainDataList = {positiveTrainPath: 1, negativeTrainPath: 0}
```

```
testDataList = {positiveTestPath: 1, negativeTestPath: 0}
```

```
h = HistogramOfGradients()
```

```
trainImages, trainDataLabels = h.findImages(trainDataList, delimiter)
```

```
testImages, testDataLabels = h.findImages(testDataList, delimiter)
```

```
trainDataIn = np.array((h.hog(trainImages)))
```

```
testDataIn = np.array((h.hog(testImages)))
```

```
NN = NeuralNetwork()  
NN.trainNN(trainImages, trainDataIn, trainDataLabels)  
NN.testNN(testImages, testDataIn, testDataLabels)
```